

Chapter 6: Conclusions, Recommendations and Future Works

6.1 Introduction

This chapter closes the chain from the problem stated in Chapter 1 to the evidence reported in Chapter 5. The problem was that small language models that run privately on consumer devices know far less than large cloud models, and that a frozen small model cannot be retrained cheaply to close the difference. The study built CAMR, an external memory engine that sits beside a frozen small model. It measured how much of the capability gap the engine closes, for which kinds of task, at what token and time cost, and for which model sizes. It then packaged the engine as an application. Section 6.2 states what the evidence established, section 6.3 what should be done because of it, and section 6.4 what remains open. No new results are introduced; each statement refers to the section of Chapter 5 that supports it.

6.2 Conclusions

The conclusions answer the research questions of section 1.4 in order.

6.2.1 Research Question i: Techniques and Their Limits On-Device

The literature review (Chapter 2) identified similarity retrieval, recency and importance scoring, compression and graph-based association as the main memory techniques, all designed for large models. Implementing them for a small model on a CPU showed which of them transfer. Recency and importance (Equation 4.1) lowered full-support recall on both multi-hop benchmarks, because benchmark questions are independent and recency carries no signal about them (section 5.8.1). Two capability-adaptive additions transferred well. Gating kept irrelevant text away from a weak reader. Entity bridging, a one-step associative hop that needs no graph construction, raised full-support recall on 2WikiMultiHopQA from 35.8% to 78.4% at the same token budget. The applicable techniques on-device are therefore cheap, similarity-based and selective, not the weighted composites of large-model agent memories.

6.2.2 Research Question ii: How Much of the Gap a Frozen-Weights Memory Engine Closes

Against the frontier cloud model (Kimi K3), the engine closed the entire single-hop knowledge gap for a 0.5B model (gap closed 1.00) and half of the multi-hop gap (0.50, 95% CI 0.20–0.90). The multi-hop result was confirmed on 200 held-out questions with the configuration fixed in advance. With CAMR, a 3B model reached 54.0% exact match against the cloud model's closed-book 51.8%, at 5.8 s instead of 18.4 s per answer, and the gain from

memory excluded zero for every small model tested (section 5.8.8). Across 19 models, 18 reached or exceeded the cloud model on long-tail facts once they had memory (section 5.8.7).

6.2.3 Research Question iii: Knowledge-Bound Versus Reasoning-Bound Tasks

Yes. Memory helps knowledge-bound tasks and does not help reasoning-bound ones. The residual gap between a 0.5B and a 20B model reading the same notes was 10.0 points on single-hop facts, 23.3 on multi-hop questions and 83.3 on mathematical reasoning (section 5.8.2). Most of the remaining multi-hop gap is a reading gap: the small model had the answer in its context for 76% of its misses (section 5.8.3). Worked-example memory for mathematics helped some model families and hurt others, depending on family rather than size.

6.2.4 Research Question iv: The Token Budget That Maximises Gap Closed Per Token

For a 0.5B model the best budget was about 128 tokens for single facts and 512 for multi-hop questions. Larger budgets lowered accuracy because extra notes distracted the small model, and reasoning tasks did not benefit at any budget (section 5.8.5). Retrieval recall saturated near 1,024 tokens (section 5.8.1). The best budget therefore depends on the task and on the model, which is why a learned policy was needed. That policy matched the cloud model's accuracy (65.3% against 63.9%) while 71% of questions stayed on the device and mean latency fell four-fold; this difference in accuracy is within sampling noise (section 5.8.9).

6.2.5 Research Question v: Cost in Tokens and Latency on a Consumer Device

The engine added a median of 187 ms and about 536 prompt tokens per answer for a 0.5B model. Peak memory was 712 MB for the engine and 1,987 MB for the model runtime, and the store was 47.8 MB (section 5.8.10). The ten-per-cent retrieval-latency target (NFR-02) was not met for the 0.5B model, because that model answers in about 180 ms. The overhead was diagnosed as CPU contention rather than engine cost: about 29 ms in isolation. Decode speed fell by 3–17% with memory, the cost of reading a longer prompt. On this class of device, models of 7B parameters or more with memory answered more slowly than the cloud model, which makes **1.5–4B the model size that suits the engine** on a CPU-only laptop.

6.2.6 Contribution and the Boundaries of the Findings

The study contributes four things: (1) a working, open-source memory engine and evaluation harness; (2) measured evidence that external memory lets small on-device models match a frontier cloud model on knowledge-bound questions, together with the residual gap that memory cannot close; (3) a size map of 19 small models showing where memory pays off;

and (4) a released desktop application, CAMR Personal, that puts the approach in users' hands. These conclusions hold for English benchmark questions, for frozen quantised models run with greedy decoding on 4-core CPUs with 16 GB of memory, for the sample sizes reported (30–200 questions per condition, 500 for retrieval), and for a stochastic cloud reference sampled once. They do not yet extend to real users' personal data, to phones, or to sessioned conversations where recency matters. Table 6.1 summarises the alignment from objectives to conclusions.

Table 6.1: Alignment from objectives to evidence and conclusions

Objective (section 1.3.2)	Implementation	Evidence (Chapter 5)	Conclusion
i. Investigate capability gaps and memory techniques	Literature review; engine built with every proposed technique switchable	5.8.1 retrieval study; 5.8.2 residual gap	Similarity + gating + bridging transfer; Eq. 4.1 does not; the gap is mostly knowledge and reading
ii. Design a lightweight memory engine	Write path, screener, single-file store, adaptive read path, budgeter (5.4)	5.5 test cases TC-01–TC-22; defects D-01–D-09	Design implemented and verified; deviations justified
iii. Develop it on a consumer laptop without a GPU	SQLite + sqlite-vec, BGE on CPU, Ollama	5.2; 5.8.10 device profile	Runs in under 3 GB; NFR-02 not met for 0.5B (contention)
iv. Evaluate gap closed by task type, with tokens and latency	Floor/treatment/ceiling/ceiling_rag harness; 19-model sweep; n = 200	5.8.2, 5.8.7, 5.8.8	Facts: fully closed; multi-hop: half, and matched by 3B; reasoning: not closed
v. Determine the best token budget	Budget sweep; learned policy	5.8.5, 5.8.9	128 tokens (facts), 512 (multi-hop) for 0.5B; learned routing matches the cloud with 71% on-device

6.3 Recommendations

The recommendations are ordered by priority. The first three are essential safeguards for anyone deploying the approach; the rest are improvements. Each one names its audience and the evidence behind it.

Table 6.2: Recommendations

#	Audience	Recommended action	Evidence and condition
1	Developers deploying on-device assistants	Pair a 1.5–4B model with external memory for personal and long-tail facts; do not rely on the small model's own knowledge	5.8.7: facts alone $\leq 33\%$, with memory 70–80%; 1.5–4B faster than the cloud. Holds on CPU-only devices
2	Developers deploying on-device assistants	Route reasoning-heavy questions (e.g. arithmetic) to a stronger model, with the user's consent, instead of adding memory	5.8.2: residual gap 83 points on GSM8K; 5.8.9: the learned router sent all maths to the cloud
3	CAMR maintainers	Code-sign the Windows and macOS downloads before wider distribution, and keep feedback strictly opt-in	5.6: unsigned builds trigger OS warnings; privacy is a stated design requirement (NFR-06)
4	Designers of memory	Use similarity ranking with gating and	5.8.1: Eq. 4.1 lowered recall; bridging

#	Audience	Recommended action	Evidence and condition
	engines for small models	entity bridging; avoid recency/importance weighting for independent questions; calibrate the abstain threshold per embedder	doubled 2Wiki recall; 5.4.7 calibration rule
5	Designers of memory engines for small models	Keep the memory budget small for tiny models (≈ 128 tokens for facts) and let a policy adapt it	5.8.5 budget curve; 5.8.9 learned policy
6	CAMR maintainers	Partition CPU threads between the model and the embedder on 4-core devices	5.8.10: p90 retrieval 157 $\rightarrow$ 50 ms with partitioning
7	Researchers evaluating memory for small models	Report the residual gap against a model reading the same notes, use a frontier ceiling, and use at least 200 questions for headline claims	5.8.2 (gap closed exceeds 1 against a weak ceiling); 5.8.8 (pilot vs $n = 200$)
8	The student and supervisor	Obtain ethics approval and a consent statement before collecting feedback from real users of CAMR Personal	5.5.3: no user study yet; feedback export is designed to be metrics-only by default

6.4 Future Works

The directions below extend the research. They do not repair promised work: everything in the approved scope was implemented and evaluated. They are ordered by expected value, and each follows from a limitation or open question in section 5.8.

Table 6.3: Prioritised future works

Priority	Direction and rationale	What changes / resources	How to evaluate
1	Longitudinal user study of CAMR Personal. Whether the assistant gets more useful as a person's own memory grows is the application's central claim, and no user has tested it yet (5.8.12).	Real users and their own notes over 4–8 weeks; ethics approval; the existing opt-in metrics export	Weekly share of answers drawn from memory,  rate, and accuracy on each user's repeated questions (day 1 vs week 4)
2	Sessioned memory benchmark. Recency could not show a benefit on independent questions (5.8.1); conversations are not independent.	A sessioned benchmark such as LongMemEval or LoCoMo; the recency term already implemented	Gap closed with and without recency on the same sessions
3	Closing the reading gap. 76% of multi-hop misses had the answer in context (5.8.3).	Structured notes, answer-extraction prompts, or a small reader-tuned model, keeping the engine fixed	Correct-when-in-context on HotpotQA and 2Wiki, $n \geq 200$
4	CAMR on phones. The Android prototype compiles and passes its unit tests but has not run on a device (5.8.11).	Two to three mid-range Android phones; on-device model files; gating calibration for the phone embedder	Retrieval share of latency (NFR-02), peak memory, battery per 100 answers, and accuracy against the desktop results
5	Stronger evaluation basis. Single-hop and reasoning results rest on pilot samples, and the cloud reference was sampled once (5.8.12).	$n \geq 200$ per task; repeated cloud samples or a deterministic open frontier model; non-English data	Confidence intervals that exclude the pilot's uncertainty; variance across repeats
6	Better routing signals. The learned router reached 61–65% against an oracle of 79% (5.8.9).	The small model's own confidence and the grounding signal as router features	Held-out accuracy and cloud share against the same oracle

Each direction should be scoped with the supervisor. The first two are the most valuable because they test the claim the application makes to its users, that it grows with them, under the conditions in which it will be used.